

CIS 455/555: Internet and Web Systems

Fall 2020

Team Project Specifications

Final code and

<https://eduassistpro.github.io/>

(last week)

1 Overview

Assignment Project Exam Help

For the term project, you will build a distributed Web indexer/crawler and analyze its performance. This will involve several components, each of which is loosely coupled.

- Crawler
- Indexer/TF-IDF Retrieval Engine
- PageRank
- Search Engine
- Experimental

[Add WeChat edu_assist_pro](https://eduassistpro.github.io/)

<https://eduassistpro.github.io/>

In addition, you may need to build a command-line interface to launch the components. More details on each component are provided in the assignment specification, which is relatively open-ended and includes many possibilities for extra credit; the EC item is subject to suggestions, and you should feel free to provide other or additional features if you like. However, you are strongly encouraged to get the basic functionality working first!!

Suggested approach: Spend some time early coordinating with your group-mates and deciding which modules from your previous homework assignments are “best of breed.” Designate one person to be responsible for each task.

Make sure adequate time is spent defining **interfaces between components** (in this case, appropriate interfaces might be REST-style messages analogous to HW3, Web service calls, and perhaps common index structures), and also plan to spend significant time integrating, **especially with EC2 and real data**. We strongly recommend that you do integration testing *regularly*, ideally every few days. Not only is this important to catch issues with compatibility, it is also helpful catch search result quality issues as well.

As in previous projects, please consider the use of automated tools for building your project (Maven) and for version control (Git, with a **private repository on BitBucket/GitHub/GitLab**).

Note that the report includes a non-trivial evaluation component.

You may want to make use of (1) Vagrant for debugging and development, (2) Bitbucket for sharing your work, (3) JUnit tests for validating that the code works (or remains working), (4) Amazon EC2 or Google Cloud to deploy and evaluate your system.

2 Project specifications

2.1 Crawler

The Web crawler should build upon your past homework assignments, and it should be able to parse typical HTML documents. It should check for and respect the restrictions in `robots.txt` and be well-behaved in terms of concurrently requesting pages. The crawler should be distributed, Mercator-style, across multiple machines. It should be an improvement over the implementation from HW3 and your basic implementation.

Just as in HW2, the crawler should track visited pages and not index a page more than once. It also *must* identify itself as 'cis455crawler' in the User-Agent header!!

Recommendations: Try to start crawling early; your team will need at least a small corpus to test the other components! It is useful to keep much of the crawler's state (URLs, visited pages, etc.) so that you can interrupt a crawl when errors happen and continue it later on. Since you will keep the crawled documents around in some distributed storage system, make sure that you can retrieve everything if something goes wrong. Please be careful not to DoS attack any

For best results, you should start crawling in a timely fashion. Remember that your team's indexing is, garbage in => garbage out. Try to avoid spider traps or starts crawling in a timely fashion. Remember that your team's indexing is, garbage in => garbage out. Try to avoid spider traps or starts crawling in a timely fashion. Remember that your team's indexing is, garbage in => garbage out.

Extra credit: Support crawling and indexing additional content types, Word documents, etc. You can also provide a way to search for images, e.g., based on the anchor tags of links that point to them, the 'alt' text that is provided in the img tags, or words that may appear in their URLs. You may use Apache Tika or PDFBox as libraries for reading these file formats.

Extra credit: Extend your crawler to collect additional off-page features about the pages it crawls, e.g., the location of the server or the age/stability of the domain registration, and provide a way to use this information for ranking.

2.2 Indexer

The indexer should take words and other information from the crawler and create a **lexicon**, **inverted index**, and any other data structures that are necessary for answering queries. Your indexer should provide a way to return weighted answers that make use of TF/IDF, proximity, and any other ranking features that are appropriate.

It should use MapReduce or Storm/StormLite to generate the data structures from the output of the crawler, either using the **StormLite framework from HW3 or Apache Hadoop/Elastic MapReduce**. It should store the resulting index data persistently across multiple nodes using BerkeleyDB or a cloud storage system like S3, DynamoDB, or RDS so lookups are fast.

Extra credit: Include document and word metadata that might be useful in creating improved rankings (e.g., the context of the words -- show a small excerpt of the original document on your results page in which the hits are highlighted, e.g., in bold).

Extra credit: Fault tolerance. Your search engine should continue to work if some of the nodes crash or are terminated. This requires at least some degree of replication (so that parts of the index do not become unavailable when crashes occur) and a way to monitor the 'health' of the nodes in the index.

Given information from crawling, you should perform **link analysis** using the PageRank algorithm, as discussed in class and in the book. You should implement this as a MapReduce job.

Extra credit: Adapt your HW3 solution to support more efficient computation across multiple iterations. Since PageRank is an iterative algorithm, you should automatically run multiple instances of a job (otherwise you'll have to trigger each iteration manually in the web interface). However, you should also consider mechanisms for reusing results across iterations. Apache Spark (not Spark Java), discussed below, may serve as inspiration. There is also a research paper on “Haloop” that might give some ideas.

This component is fairly self-explanatory, as the goal is to provide a search form and weighted results list. One aspect that will take some experimentation is determining how to combine the various factors (PageRank, TF/IDF, other word features) into a single ranking of the results. This user interface portion of

this must make use of route handlers in Spark Framework or your HW1 server. You may leverage open-source client-side libraries such as Vue.js, Bootstrap, etc.

Recommendations: Ranking performance is crucial - the best scalable/secure/robust search engine design doesn't help if your search results are not useful! Therefore, you need to spend at least a few days on tuning your ranking function. This is challenging because the ranking component depends on all the other components, so you should insist that your teammates quickly provide simple but functional prototypes of their components (single-node crawler, basic indexer, basic PageRank, etc.), so that you can start testing.

It is also a good idea to have a 'debug mode' in which your engine shows additional information about how the results were ranked, and why (e.g., the raw PageRank scores, the TF and IDF values, and so on), so you can do example queries ("A te why the results you expected do not show up at o

Extra credit: Integrate search results from Amazon or eBay, weather forecasts, business info from Yelp, results from Facebook, etc. Simple implementations could display these results in a sidebar or in a box above or below the other results; more advanced implementations could merge the external results with the search engine's own results (but please visibly indicate which is which!).

Extra credit: Implement a simple Google-style spell-check for word (e.g., adding, removing, transposing characters) and see if simple edits to the word is "nearby."

Extra credit: Consider an autocomplete feature which entries are "goo le example of this would be provide feedback about

3 Experimental Analysis and Final Report

Building a Web system is clearly a very important and ch important is being able to convince others (your managers, instructors, peers) that you succeeded. We would like you to actually *evaluate* the performance of your methods, for instance relative to scalability.

For evaluation, you should log into multiple Amazon EC2 nodes and run the system.

One approach is to use the system with one, two, up to n nodes (where n is, say, 10 EC2 nodes), and compare overall performance, e.g., crawl throughput, time to run PageRank or to build the index, time to answer queries, etc. The result of such an experiment could be graphs that show some performance metric (e.g., crawl throughput) versus the number of nodes. For benchmarking query performance, you can write a simple query generating tool that submits many queries at once, and compare response time. What is the maximum number of concurrent requests you can reasonably handle, when varying the number of nodes? Can you separate out the overhead of the different components (including network traffic)? What is the bottleneck? Etc.

Your final report (a PDF document of six pages or less, due on the day of your demonstration) should include at least:

- Introduction: project goals, high-level approach, milestones, and division of labor

- project plan
first two sect
architecture, (2) some rough milestones, (3) your private bitbucket repository (for which you should grant read permissions to the email upercis455@gmail.com), and (4) a list of team members with a **division of labor**. Create your project group on Canvas and submit the project plan to the group.
- To share the bitbucket repo, go to the repository settings => "Repository access" => "Add group access". Then add the group `edu_assist_pro` and click "Add".

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

a. To share the bitbucket repo, go to the repository settings => "Repository settings" => "Access and group access" Then add click "Add"

a. To share the bitbucket repo, go to the repository settings => "Repository settings" => "User and group access". Then add assist@cs.gatech.edu click "Add"

- b. To create your project group, go to Canvas and click the “People” page => “Groups” tab. On this tab, you should be able to add yourself to a group. Note that some groups are 4-person and some are 5-person. **Add yourself to a group before submitting to ensure that everything is tracked correctly!**

3. Your “code-complete” code submission, **due in the Bitbucket repository you list in your report.**

The final version of the code submission, due with your demonstration, must be in the repository at that point. It must contain a) the entire source code for the project, as well as any supplementary files needed to build your solution, b) a README file, and c) the final report in point #5 below. The README file must contain 1) the full names and SEAS login names of all the project members, 2) a description of all features implemented, 3) any extra credit claimed, 4) a list of source files included, and 5) detailed instructions on how to install and run the project. The code must contain a reasonable amount of documentation. Please make this repository **private** and then share it with us.

4. Each team will need to **schedule a project demo** during the final exam period (or earlier if you prefer). All team members must be able to access the code using your search engine on Amazon. <https://eduassistpro.github.io/>
5. The **final report**, due by the end of the exam period, must include all the information from the project plan (possibly revised and/or with more details), plus a description of your implementation, your evaluation results, and your conclusions. **The final report should be submitted by one member of your team via Canvas.** f six pages or less. It

You may use the standard Java libraries, code from previous CI assignments by you or your team members, and any code we provide or explicitly approved by us. Please don't be tempted to reuse code from previous assignments by you or your team members. [Add WeChat edu_assist_pro](https://eduassistpro.github.io/)

To maintain fairness (a Piazza post for approval), we will create a Piazza post for approval along with a short description of the project. We will then post a follow-up that indicate you should post it there, and we will then post a follow-up that indicate

5 Hints

Based on last year's project, many teams tend to divide up the work at the beginning and then meet again a few days before the deadline. This rarely works well; most people tend to under-estimate the amount of work that is needed to integrate the various components at the end. To avoid this, we recommend that you:

- Define clear interfaces at the beginning. Ideally, write a few unit tests, so that everyone understands what their component is supposed to do, and how it interacts with the other components.
- Do integration testing as early as possible. Ideally, everyone builds a very simple demo version of their component first, exchanges that code with the others, and then adds features one by one.
- Meet regularly and keep everyone posted on your progress.
- **Start early!!!**